

Toward Statistical Machine Translation without Parallel Corpora

Abstract

Current statistical machine translation methods rely on very large parallel corpora to achieve state of the art performance. However, such resources are only available for very few language pairs as they are very expensive to obtain. On the other hand, monolingual data (such as newswire) is widely available and cheap to collect. In this work, we propose a set of cues derived from monolingual resources and systematically introduce them in a phrase-based machine translation pipeline *in place* of features induced from parallel data. We show on two language pairs that monolingual data can take us a long way toward inducing a high quality machine translation system.

1 Introduction

Current statistical machine translation (SMT) methods (e.g. (Koehn et al., 2003; Chiang, 2005)) crucially rely on vast amounts of sentence aligned translations in order to achieve state of the art performance. These resources are only available for very few language pairs because producing them in sufficient quantities is an expensive and time consuming endeavor. Moreover, the SMT system performance tends to drop if test data comes from a different domain than the parallel data used in training. One general idea to deal with data sparsity is to attempt to collect (more) parallel data automatically (e.g. (Munteanu and Marcu, 2006; Smith et al., 2010; Uszkoreit et al., 2010)). However, some assumptions are typically made about the comparable bilingual corpus (such as document level alignment) used for mining near parallel text fragments.

In this work, we approach the problem from an entirely different perspective: we use monolingual resources directly to induce an end-to-end statistical machine translation system. In particular, we extend a long line of work on inducing translation lexicons (e.g. (Rapp, 1995; Fung and Yee, 1998; Koehn and Knight, 2000; Klementiev and Roth, 2006; Haghighi et al., 2008; Mimno et al., 2009)) to induce translation features, and propose a novel algorithm for estimating reordering features for the phrase-based machine translation framework (Koehn et al., 2003). Much of the prior work on lexicon induction is motivated by the resource constrained SMT, however, to the best of our knowledge, this work is the first attempt to extend and apply these ideas in an end-to-end machine translation pipeline.

The rest of the paper is organized as follows. Section 2 begins with the relevant background on the phrase-based SMT framework we will use in the rest of the paper and continues to give a brief overview of the existing work on inducing translation lexicons. Section 3 motivates and introduces translation and reordering features induced from monolingual data alone. Section 4 studies the informativeness of these features as they are added to the Machine Translation pipeline. Finally, Section 6 concludes and discusses directions for future work.

2 Background

2.1 Phrase-Based SMT

Review phrase-based setup a-la (Koehn et al., 2003).

Need a more concrete statement.

Need a good MT adaptation reference

Need a better "however" sentence.




Figure 1: Example alignment along with three kinds of orientations: monotone (m), swap (s), and discontinuous (d).

- Log linear formulation:

$$p(\mathbf{e}|\mathbf{f}) \propto \exp \sum_{i=1}^n \lambda_i h_i(\mathbf{e}, \mathbf{f})$$

$$\hat{e} = \arg \max_e p(\mathbf{e}|\mathbf{f})$$

- *Phrase extraction.* Size of the phrase table and maximum phrase length: show our plot of performance vs. max phrase length. Mention what we choose max phrase length 3 for our experiments.
- *Phrase features.* Phrase translation probability and lexical translation probability.
- *Lexicalized reordering features.* See Figure 1.
- *Other features.* Language model, penalties.

In this work, we will use the same general mathematical formulation, but propose alternative features derived directly from monolingual data.

2.2 Lexicon Induction

Let us now briefly review relevant prior work on lexicon induction. In Section 3 we will build on this this

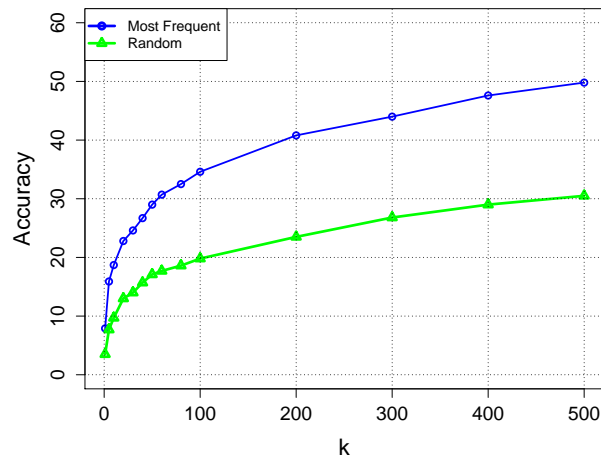


Figure 2: Accuracy on wikipedia for most frequent and random 1000 source words.

work specifically to propose alternatives to some of the types of features we outlined in Section 2.1.

TODO: Alex: Fill in.

Most of the previous work evaluates results on a small set of hand selected words (e.g. 100 nouns in (Rapp, 1995)). However, if the objective is to induce large translation tables, as it is in our work, the reported results can be misleading. **TODO: Describe**

Figure 2.

Diss prev work, but make sure that it comes across that these features are informative.

3 Reducing the Parallel Data requirement / Estimating Parameters from Monolingual Data

3.1 Phrase extraction

Building on the many successful efforts in bilingual lexicon induction from monolingual corpora, we compile a table of *phrasal* translations from monolingual corpora. Current methods for inducing a bilingual lexicon, or, equivalently, a unigram phrase translation table, are computationally expensive as each source and target word pair must be scored. That is, the models must compute $|V_s| * |V_t|$ similarity scores, where $|V_s|$ and $|V_t|$ are the sizes of the source and target vocabularies, respectively. As we search for longer phrase pairs, this search space increases exponentially with the size of the n -grams. Exhaustively scoring all phrase pairs up to length

three requires $|V_s|^3 * |V_t|^3$ score computations. We prune the phrase pair search space using methods from Information Retrieval, like that used in Uszko-reit et al. (2010).

First, we compose a set of all n -grams up to length three in the source side monolingual corpus and in the target side monolingual corpus. We store each source side n -gram’s frequency along with the phrase. We also store each target side n -gram in a frequency inverted index. We index target side phrases by frequencies within a band of their actual observed frequency in the monolingual corpus. Additionally, for the target side n -grams, we look up each word in the phrase in our bilingual dictionary, and we store a set of source side words that translate into any word in the phrase.

These sets of observed phrases, frequencies, and the bilingual dictionary together allow us to effectively filter our search space using combinations of inverted indices. Given a source side test set, we collect n -grams up to length three (in our Urdu test set, there are 36,423 phrases), as we did for the monolingual corpora. We then look up each source side phrase’s monolingual corpus frequency and consult the target side inverted index to find a list of target side phrases that occur in the same frequency band. Of those target side phrases, we keep only the ones with at least one word translating into a word in the source side phrase.

This method of pruning the phrase pair search space involves two manually tuned parameters. In order to evaluate the filtered phrase tables and tune the parameters, we used the Moses decoder’s trace function to find the set of phrases used in decoding an Urdu test set. We compare our filtered phrase tables to this set of phrase translation rules and attempt to maintain as many of them as possible, while pruning the set of phrase pairs down to manageable size.

Our baseline phrase table is generated using a bilingual dictionary. For each Urdu test set phrase up to length three, we generated English phrases from all combinations of dictionary translations and all possible reorderings. For the baseline and our pruning methods, the number of filtered phrase pairs and the percent of phrases used by the Moses decoder not pruned away are given in Table 3.1.

Second round of pruning: after monolingual feature extraction, before re-ordering estimation.

Needs to be discussed after explanation of those methods?

3.2 Phrase scoring

In place of phrase translation probabilities estimated from bilingual alignments, we propose to compute similarity scores computed (almost) solely from monolingual resources.

Contextual similarity. We extend the vector space approach of (Rapp, 1999) to compute similarity between *phrases* in source and target language. More formally, assume that (f_1, f_2, \dots, f_N) and (e_1, e_2, \dots, e_M) are (arbitrarily indexed) source and target vocabularies, respectively. A source phrase f (target phrase e) is represented with an N (M) dimensional vector. Only the components corresponding to words that appear in the context of f (e) in data take on non-zero values, which typically measure how “unique” a word is to the context in the dataset. Next, f ’s contextual vector is projected by mapping each component to a component in the target space corresponding to its translation (taken from a small seed dictionary), but retaining the source component value. Finally, the pair (f, e) is scored by computing similarity between the (projected) source and target vectors. Various means of computing the component values and vector similarity measures have been proposed in literature (e.g. (Rapp, 1999; Fung and Yee, 1998)). While the quality of the resulting induced lexicon depends on the data, we found the following to work best in our experiments. We compute the value of the k -th component of f ’s contextual vector as follows:

$$w_k^{(i)} = n_{i,k} \times (\log(n/n_i) + 1)$$

where $n_{i,k}$ and f_k are the number of times f_k appears in the context of f_i and in the entire corpus, and n is the maximum number of occurrences of any word in the data. Intuitively, the more frequently f_k appears with f_i and the less common it is in the corpus in general, the higher its component value. Similarity between two resulting vectors is measured as a cosine of the angle between them.

Temporal similarity. Online content is often published along with temporal information: news

discuss dictionary(ies)

discuss dictionary parameters

data details?

Pruning filters	Phrase Pairs	Percent of total search space	Findable types	Findable tokens
Unpruned phrase table	37,322,465,985	100%	100%	100%
Baseline phrase pairs	29245036	0.08%	15.37	25.07
Frequency-based pruning	4,450,429,494	11.92%	85.21%	87.48%
Frequency and Dictionary pruning	1,436,823,109	3.85%	57.79%	58.76%

Table 1: This shows the tradeoff between pruning the phrase pair search space and the accuracy of the final set of phrase pairs. The findable types and tokens measures refer to the percent of phrase types and tokens used by Moses to decode a test set that are not pruned away.

feeds, for example, are comprised of news stories annotated with date and time of publication. The feeds are specialized for the target geographical locations and vary in content across languages. Still, many events are deemed relevant to multiple audiences and the news stories related to them appear in several languages, although rarely as direct translations of one another. Phrases associated with these events will appear with increased frequency in multiple languages around the dates when these events are reported. Such weak synchronicity provides a cue about the relatedness of phrases across the two languages. In order to score a pair of phrases across languages, we can compute the similarity of their temporal signatures. To generate a time sequence for a given word, we first sort the set of (time-stamped) documents of our corpus into a sequence of equally sized temporal bins. We then count the number of occurrences of a phrase in each bin. Changing the size of the bin or computing counts in a sliding window instead can recover some accuracy if the temporal alignment between two languages in our dataset is poor (Klementiev and Roth, 2006). Finally, we normalize the sequence and use either the cosine measure to score similarity.

Orthographic / phonetic similarity. Etymologically related words often retain similar spelling across languages with the same writing system, and the edit distance can be used to measure their orthographic similarity. We extend this idea to phrases by using word alignments within a phrase pair (see Section 3.1): we score pairs of aligned words and normalize by their average length.

We can further extend this idea to language pairs not sharing the same writing system, since many cognates and transliterated words are phonetically similar. Following (Virga and Khudanpur, 2003;

Irvine et al., 2010), we treat transliteration as a monotone character translation task and use a generative model to propose a transliteration of tokens in a source phrase. Once the source words are mapped to the target writing system, the phrase similarity is computed as before.

Depending on the available monolingual data (and its associated metadata), various other similarity scores can be computed and added to the list (see, e.g. (Schafer and Yarowsky, 2002)).

3.3 Reordering

Describe the algorithm (Figure 3) for estimating orientation probabilities¹. Talk about the issue of too much weight on out-of-order orientation. Talk about where w comes from. Mention that it is not as expensive as it looks. The intuition is simple - see if translated phrases tend to keep the order.

4 Experiments

4.1 Data

Describe data we use in the experiments: Europarl (Koehn, 2005), Gigaword, our own crawls².

4.2 Single language

1. *Phrase features.* (a) Augment phrase scores with mono features. If we see better performance, reduce the amount of parallel data until it matches the performance of the original system. Make the tradeoff argument. (b) **(lesion experiments)** See how well we do with mono features alone.
2. *Orientation features.* Use mono orientation features.

¹ $\#_S(x)$ returns the count of object x in multiset S .

²Promise to distribute after publication.

Argue that enough training data is easy to get

Make sure it is correct.

Input: Source and target phrases f and e , source and target corpora C_f and C_e , phrase table pairs $T = \{(f^{(i)}, e^{(i)})\}_{i=1}^N$.
Output: Orientation features (p_m, p_s, p_d) .

$S_f \leftarrow$ sentences containing f in C_f ;
 $S_e \leftarrow$ sentences containing e in C_e ;
 $(B_f, -, -) \leftarrow \text{CollectOccurs}(f, \cup_{i=1}^N f^{(i)}, S_f)$;
 $(B_e, A_e, D_e) \leftarrow \text{CollectOccurs}(e, \cup_{i=1}^N e^{(i)}, S_e)$;
 $c_m = c_s = c_d = 0$;

foreach $f_b \in B_f$ **do**
 foreach translation e^* of f_b in T **do**
 $c_m = c_m + \#_{B_e}(e^*)$;
 $c_s = c_s + \#_{A_e}(e^*)$;
 $c_d = c_d + \#_{D_e}(e^*)$;

$c \leftarrow c_m + c_s + c_d$;
return $(\frac{c_m}{c}, \frac{c_s}{c}, \frac{c_d}{c})$

CollectOccurs(r, R, S)
 $B \leftarrow ()$; $A \leftarrow ()$; $D \leftarrow ()$;
 foreach sentence $s \in S$ **do**
 foreach occurrence of phrase r in s **do**
 $B \leftarrow B + (\text{longest preceding and in } R)$;
 $A \leftarrow A + (\text{longest following and in } R)$;
 $D \leftarrow D + (\text{longest discontinuous and in } R)$;
 return (B, A, D) ;

Figure 3: Estimating reordering probabilities from monolingual data.

3. Induce phrase table.

4. *Put everything together.* Run the entire pipeline.

4.3 Big experiment

Now, run the entire pipeline on a handful of languages extracting monolingual features from the Gigaword and our crawls.

5 Discussion

6 Conclusions and Future Work

First to make use of plentiful monolingual data to reduce the dependence on expensive parallel data. In particular:

- Showed that augmenting standard pipeline with monolingual features helps.
- Demonstrated that monolingual features are informative enough on their own for a competitive system.

- Proposed an algorithm for estimating orientation probabilities from monolingual data alone.
- Build complete systems for X low-resource languages.

References

- David Chiang. 2005. A hierarchical phrase-based model for statistical machine translation. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- Pascale Fung and Lo Yuen Yee. 1998. An ir approach for translating new words from nonparallel, comparable texts. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 414–420.
- Aria Haghighi, Percy Liang, Taylor Berg-Kirkpatrick, and Dan Klein. 2008. Learning bilingual lexicons from monolingual corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- Ann Irvine, Chris Callison-Burch, and Alexandre Klementiev. 2010. Transliterating from all languages. In *Proc. of the Conference of the Association for Machine Translation in the Americas (AMTA)*.
- Alexandre Klementiev and Dan Roth. 2006. Weakly supervised named entity transliteration and discovery from multilingual comparable corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- Philipp Koehn and Kevin Knight. 2000. Estimating word translation probabilities from unrelated monolingual corpora using the EM algorithm. In *Proc. of the National Conference on Artificial Intelligence (AAAI)*.
- Philipp Koehn, Franz Josef Och, and Daniel Marcu. 2003. Statistical phrase-based translation. In *Proc. of the Annual Meeting of the North American Association of Computational Linguistics (NAACL)*, pages 48–54.
- Philipp Koehn. 2005. Europarl: A parallel corpus for statistical machine translation. In *Proc. of the Machine Translation Summit*.
- Adam Lopez. 2008. Statistical machine translation. *ACM Computing Surveys*, 40(3):1–49, August.
- David Mimno, Hanna Wallach, Jason Naradowsky, David Smith, and Andrew McCallum. 2009. Polylingual topic models. In *Proc. of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Dragos Stefan Munteanu and Daniel Marcu. 2006. Extracting parallel sub-sentential fragments from non-parallel corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 81–88.

- Reinhard Rapp. 1995. Identifying word translations in non-parallel texts. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 320–322.
- Reinhard Rapp. 1999. Automatic identification of word translations from unrelated english and german corpora. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 519–526.
- Charles Schafer and David Yarowsky. 2002. Inducing translation lexicons via diverse similarity measures and bridge languages. In *Proc. of the Annual Conference on Computational Natural Language Learning (CoNLL)*, pages 146–152.
- Jason R. Smith, Chris Quirk, and Kristina Toutanova. 2010. Extracting parallel sentences from comparable corpora using document level alignment. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 403–411.
- Jakob Uszkoreit, Jay M. Ponte, Ashok C. Popat, and Moshe Dubiner. 2010. Large scale parallel document mining for machine translation. In *Proc. of the International Conference on Computational Linguistics (COLING)*.
- Paola Virga and Sanjeev Khudanpur. 2003. Transliteration of proper names in cross-lingual information retrieval. In *Proc. of the Annual Meeting of the Association for Computational Linguistics (ACL)*.